

## Goals

Prediction in general

Plug-in approximation of prior predictive distributions---the case of the gamma-Poisson model

Plug-in approximation of posterior predictive distributions---the case of the gamma-Poisson model

Comparing prior and posterior predictive distributions

Simulating posterior (or prior) predictive distributions for whole samples

# Eliciting prior and posterior predictive distributions through MC approximation

## CSSS/STAT 564: Bayesian Statistics for the Social Sciences

William Brown and Alex Jiang

CS&SS 564 (Spring 2021)

## Goals

- Illustrate plug-in approximation of DAG-based joint probability distributions and the marginal distributions they imply
- A practical guide for eliciting prior predictive distributions and posterior predictive distributions
- Predicting identically generated data sets

## Prediction in general

- Our intention in fitting probability models to data is often to understand the DGP underlying the data, often for the sake of making informed decisions about future outcomes.
- Suppose that the underlying data generating process is as follows:  $Y$  is generated according to a distribution characterized by (possibly unknown) parameter  $\theta_m$  --  $Y \sim \text{Dist}_m(\theta_m)$ . Our goal is to describe the distribution of new values  $\tilde{Y}$  coming from such distribution, before there are observed.
  - such procedure of making inferences about an unknown observable is often called **predictive inference**.
  - Let  $\tilde{Y}$  denote a "predictand"---a class of unknown quantities for which we intend to estimate future values. this quantity has the same domain as realized values of  $Y$ :  $\tilde{Y} = y \in \mathcal{Y}$ . For greater succinctness, we can abbreviate the expression  $\tilde{Y} = y$  more simply as  $\tilde{y}$ .

**Question:** how do we do predictive inference, given the model we just described?

- If  $Y \sim \text{Dist}_m(\theta_m)$  and we knew  $\theta_m$  exactly, then  $\text{Dist}_m(\theta_m)$  would double as a **predictive distribution**, i.e.  $\tilde{Y} \sim \text{Dist}_m(\theta_m)$ 
  - However, we never know the true value  $\theta_m$  (at least practically speaking). All we know (based on the model we specified) is the prior distribution of  $\theta$ .

**Bayesian prediction framework allows us to incorporate our (modeled) uncertainty regarding this fixed but unknown quantity.**

- When we don't know  $\theta_m$ , our prediction of  $\tilde{Y}$  requires us to elicit predictive probabilities  $p(\tilde{y})$  by applying the Rule of Total Probability to the joint distribution  $p(\tilde{y}, \theta_m)$ :

$$p(\tilde{y}) = \int_{\Theta_m} p(\tilde{y}, \theta_m) d\theta_m$$

- The joint probability function  $p(\tilde{y}, \theta_m)$  can be decomposed into (A) a model for the DGP underlying  $\tilde{Y}$  conditional on  $\theta_m$ , and (B) a model for our uncertainty/belief regarding  $\theta_m$ :

$$p(\tilde{y}, \theta_m) = p(\tilde{y}|\theta_m) \times p(\theta_m)$$

$$\therefore p(\tilde{y}) = \int_{\Theta_m} p(\tilde{y}|\theta_m) p(\theta_m) d\theta_m$$

- An interpretation: Ideally we would want to sample  $\tilde{y}$  from  $p(\tilde{y}|\theta_{true})$ , where  $\theta_{true}$  is the true value of the unknown parameter. Since we don't know the true value of  $\theta_m$ , we incorporate our prior belief about this parameter to 'faciliate our guess', by taking a 'weighted average' wrt the prior distribution of  $\theta_m$ .

Depending on whether we have observed data  $Y = y$  coming from the same process, the predictive distribution we use to describe the unknown data  $\tilde{y}$  can be categorized into **prior predictive distribution** and **posterior predictive distribution**.

- the marginal distribution  $p(\tilde{y})$  above is also called the **prior predictive distribution**, because:
  - (**prior**): it is not based on observed data:
  - (**predictive**): it is the distribution of  $\tilde{y}$ , which is observable (we're about to observe it) but unknown (we haven't observed it yet)

We will learn how to sample values from the prior predictive distribution through plug-in approximation in the following section:

- This implies a very simple DAG decomposition of the full probability model:


```
# Illustration of a DAG for a univariate DGP when the model parameters are unknown
theme_set(theme_dag_gray())

simpleDAG.coords = list(
  x = c(
    theta_m = 1, Y.tilde = 2
  ),
  y = c(
    theta_m = 1, Y.tilde = 1
  )
)

simpleDAG = dagify(
  Y.tilde ~ theta_m, # Child on the left of the tilde, parent(s) on the right of the tilde separated by '+'

  coords = simpleDAG.coords
)

simpleDAG %>%
  ggdag() +
  geom_dag_node(color = "dark gray", internal_color = "dark gray") +
  geom_dag_text(col = "black")
```



## Plug-in approximation of prior predictive distributions---the case of the gamma-Poisson model

- While there are a small handful of analytical solutions for marginal predictive distributions  $p(\tilde{y})$ , these are mostly for "toy" problems. In most realistic data-analytic circumstances, plug-in approximation is favored as an efficient alternative.
- Basic procedure: draw a large MC sample from the joint PD of  $\tilde{Y}$  and  $\theta_m$ , then summarize only the marginal sample distribution of  $\tilde{Y}$ .
- Recall: sampling from a joint PD model based on a DAG is straightforward:
  - Simulate  $S$  observations from the DAG's root nodes (quantities represented by nodes with no parents):  $\{\theta_m^{(1)}, \dots, \theta_m^{(S)}\} \stackrel{iid}{\sim} \text{Dist}_{\theta_m}$

- Pass this output as input for the simulation of child quantities:  $y^{(s)} \sim \text{Dist}_m(\theta_m^{(s)})$
- Repeat this process until the terminal nodes are reached (quantities represented by nodes with no children).
- Where does  $p(\theta_m)$  come from?
  - In general, it is some degree-of-belief budget about this uncertain quantity,  $p(\theta_m|\eta)$ , where  $\eta$  are the hyper-parameters of this model.
  - It could be a prior PD model,  $p(\theta_m|\eta_0)$ , where  $\eta_0$  are the prior hyper-parameters.
  - It could be a posterior PD model,  $p(\theta_m|\eta_0, y)$ , where  $\eta_0$  are again the prior hyper-parameters and  $y$  abbreviates observed data used to further condition or update our belief about  $\theta_m$ .

*# Illustration of a DAG for a univariate Poisson DGP when Lambda is unknown*

```
PriorPredDistDAG.coords = list(
  x = c(
    lambda = 1, Y.tilde = 2
  ),
  y = c(
    lambda = 1, Y.tilde = 1
  )
)

PriorPredDistDAG = dagify(

  Y.tilde ~ lambda,

  coords = PriorPredDistDAG.coords
)

PriorPredDistDAG %>%
  ggdag() +
  geom_dag_node(color = "dark gray", internal_color = "dark gray") +
  geom_dag_text(col = "black")
```

lambda

Y.tilde

$$\tilde{y} \in \{0, 1, 2, \dots\} \quad \lambda \in \mathbb{R}_{>0}$$

$$p(\tilde{y}, \lambda | a_0, b_0) = P(\tilde{y} | \lambda) \times p(\lambda | a_0, b_0)$$


$$\tilde{Y} \sim \text{Poisson}(\lambda)$$

$$\lambda \sim \text{Gamma}(a_0, b_0)$$

$a_0$  is a prior shape hyper-parameter and  $b_0$  is a prior concentration hyper-parameter.

```
# MC simulation of the joint distribution of lambda and the Poisson-distributed estimand based on a prior model of lambda
a_0 = b_0 = 1; S=100000; set.seed(564)
lambdaSim = rgamma(n = S, shape = a_0, rate = b_0)
Y.tildeSim = rpois(n = S, lambda = lambdaSim)

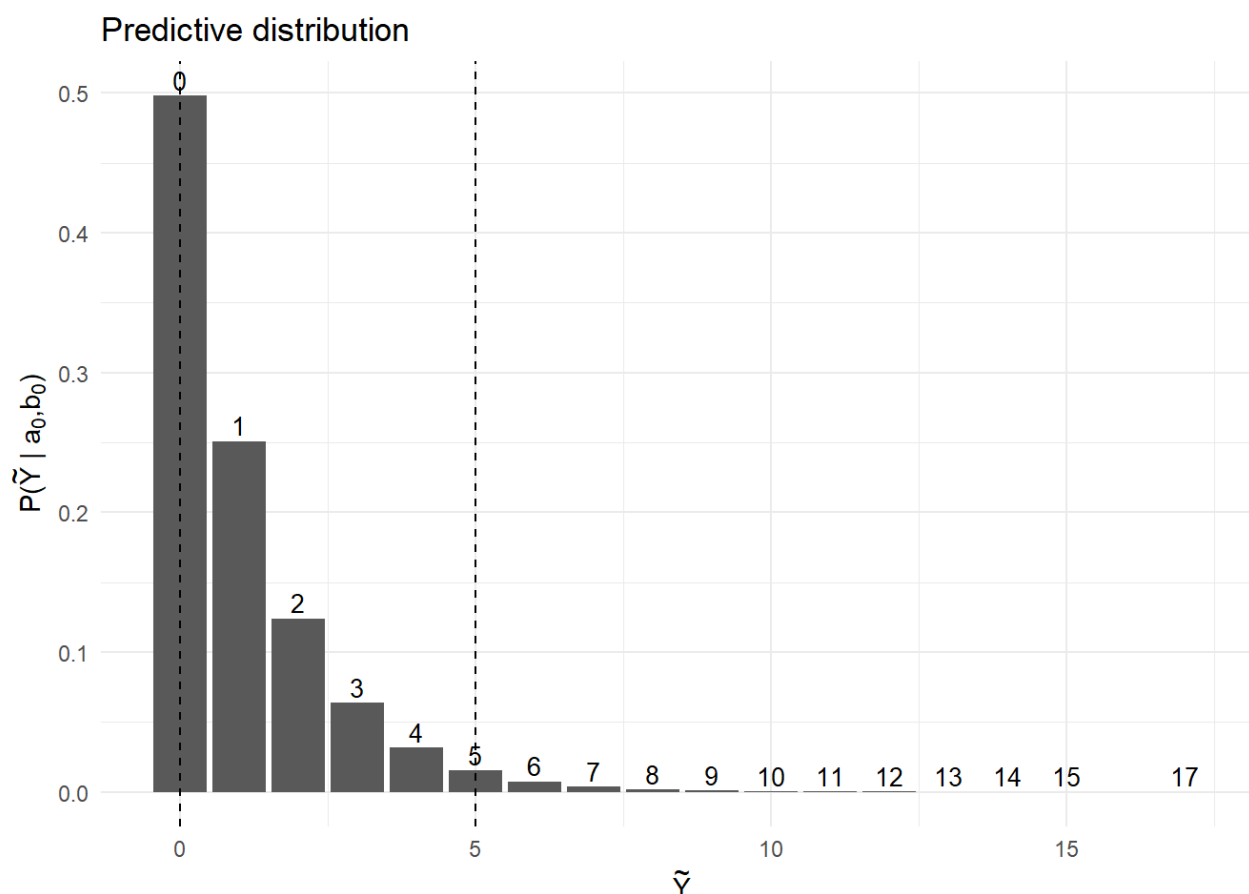
# Illustration of the approximated joint prior distribution of lambda and the Poisson-distributed estimand
jointPlugInApprox = ks::kde(x = cbind(lambdaSim, Y.tildeSim))
plot(
  x=jointPlugInApprox, # For help plotting kde objects, see ?plot.kde
  cont = c(50, 95), # Percentiles to draw highest-probability regions
  xlim = c(0,4), ylim = c(0, 6),
  xlab = expression(lambda), ylab = expression(tilde(Y))
)
```



(Note that this contour plot treats  $\hat{Y}$  as if it were continuous rather than discrete.)

- The marginal prior predictive distribution of  $\tilde{Y}$  can then be approximated by plotting out only the simulated vector  $\{\tilde{Y}^{(1)}, \dots, \tilde{Y}^{(S)}\}$ :

```
# Illustration of the approximated marginal prior distribution of the Poisson-distributed estimand
df <- tibble(x = as.integer(rownames(table(Y.tildeSim))),
             y = table(Y.tildeSim)/S)
ggplot(df, aes(x=x, y=y)) + geom_bar(stat="identity") +
  geom_text(aes(label=x), vjust=-0.3, size=3.5)+theme_minimal() +
  xlab(expression(tilde(Y))) + ylab(expression(paste("P(",tilde(Y)," | a"[0],",b"[0],
  ")")))+
  geom_vline(xintercept = quantile(x = Y.tildeSim, probs = c(0.025,.975)), lty = 2) +
  ggtitle("Predictive distribution")
```



Plug-in prior predictive distribution approximation for the gamma-Poisson model. The dashed vertical line shows the 95% quantile-based credible interval.

## Plug-in approximation of posterior predictive distributions---the case of the gamma-Poisson model

If we already observed data from the data-generating process  $y$ , the distribution of new unknown observable  $\tilde{y}$ , given the previous data  $\tilde{y}$  (note their notational difference!) is called **posterior predictive distribution**, denoted as  $p(\tilde{y}|y)$ .

- How we derive  $p(\tilde{y}|y)$  is very similar to what we did before: we first apply the Rule of Total Probability

$$p(\tilde{y} | y) = \int_{\Theta_m} p(\tilde{y}, \theta_m | y) d\theta_m$$

- We can similarly decompose the joint posterior probability function  $p(\tilde{y}, \theta_m | y)$  into (A) a model for the DGP underlying  $\tilde{Y}$  conditional on  $\theta_m$ , **given** our prior knowledge of  $y$ , and (B) the posterior distribution for our unknown parameter  $\theta_m$  given  $y$ :

$$p(\tilde{y}, \theta_m | y) = p(\tilde{y} | \theta_m, y) \times p(\theta_m | y) \\ \therefore p(\tilde{y} | y) = \int_{\Theta_m} p(\tilde{y} | \theta_m, y) p(\theta_m | y) d\theta_m$$


- Finally, based on our DAG model, we have **assumed conditional independence** of  $y$  and  $\tilde{y}$  given  $\theta$ :

$$p(\tilde{y} | y) = \int_{\Theta_m} p(\tilde{y} | \theta_m, y) p(\theta_m | y) d\theta_m = \int_{\Theta_m} p(\tilde{y} | \theta_m) p(\theta_m | y) d\theta_m$$

- Derivation of this result requires theory on d-separation for DAGs and is beyond the scope of today's lecture, but here's one possible interpretation for this:  $\theta_m$  is the common and only cause for  $\tilde{y}$  and  $y$ , so we wouldn't gain extra information about  $\tilde{y}$  if we already controlled/conditioned on  $\theta_m$ , even if we observed  $y$  in the first place.
- This is called 'Principle of the Common Cause' and you can find more information here if you're interested: <https://www.andrew.cmu.edu/user/scheines/tutor/d-sep.html> (<https://www.andrew.cmu.edu/user/scheines/tutor/d-sep.html>)
- To illustrate, let us first generate a sample of 5 observations from a known Poisson distribution:

$$\{Y_1, \dots, Y_5\} \stackrel{iid}{\sim} \text{Poisson}(\lambda = 5)$$

```
# Illustration of a known {Y_1,...,Y_5} ~ Poisson(lambda=5) sample
check_occurences <- function(x) {
  return(sum(Y.Sim %in% x))
}
Y.Sim = rpois(n = 5, lambda = 5)
df <- tibble(x = 0:20, y1 = dpois(x = 0:20, lambda = 5),
             y2 = sapply(0:20, check_occurences)/sum(Y.Sim)) %>%
  pivot_longer(!x, names_to = "type", values_to = "probs")
ggplot(df) + geom_bar(stat='identity',
                     aes(x = x, y = probs, fill = factor(type)),
                     position=position_dodge(width = .7), width = .7)+
  theme_bw() + xlab("Y") + ylab(expression(paste("P(Y | ", lambda, ") or sample proport
ions"))) +
  theme(legend.title = element_blank()) +
  scale_fill_discrete(labels = c(expression(paste("Poisson(", lambda, "=5)")), "Sample di
stribution")))
```



- Eliciting posterior predictive distributions involves simulating from a joint distribution that includes not only the predictand and model parameter(s) but also observed data:



*# Illustration of a DAG for a univariate Poisson DGP when lambda is unknown, with both observed and unobserved/predicted values of Y*


```
PostPredDistDAG.coords = list(
  x = c(
    lambda = 1, Y_1 = 2, Y_2 = 2, Y_3 = 2, Y_4 = 2, Y_5 = 2, Y.tilde = 2
  ),
  y = c(
    lambda = 1, Y_1 = 3.5, Y_2 = 2.5, Y_3 = 1.5, Y_4 = 0.5, Y_5 = -0.5, Y.tilde = -1.5
  )
)

PostPredDistDAG = dagify(

  Y_1 ~ lambda,
  Y_2 ~ lambda,
  Y_3 ~ lambda,
  Y_4 ~ lambda,
  Y_5 ~ lambda,
  Y.tilde ~ lambda,

  coords = PostPredDistDAG.coords
)

PostPredDistDAG %>%
  ggdag() +
  geom_dag_node(color = "dark gray", internal_color = "dark gray") +
  geom_dag_text(col = "black")
```



- However, in this case, we know several values represented in the joint distribution so we must apply a different strategy for sampling from the DAG. The easiest and fully sufficient way is to simulate from the posterior distribution of  $\lambda$ , then to simulate the predictand as before.
- For the gamma-Poisson model, an analytical solution is available for the posterior distribution with the following posterior shape and concentration hyper-parameters:


$$\lambda|y \sim \text{Gamma} \left( a_n = a_0 + \sum_{i=1}^n y_i, b_n = b_0 + n \right)$$

```
# MC simulation of the joint distribution of lambda and the Poisson-distributed estima
nd based on a posterior model of lambda
a_n = a_0+sum(Y.Sim); b_n = b_0+5
lambdaPostSim = rgamma(n = S, shape = a_n, rate = b_n)
Y.tildePostSim = rpois(n = S, lambda = lambdaPostSim)


# Illustration of the likelihood, prior, and posterior over Lambda
likelihoods = exp(colSums(apply(
  X = matrix(seq(0, 15, 0.01), ncol = 1),
  MARGIN = 1,
  FUN = dpois, x = Y.Sim, log = TRUE
))))

df1 <- tibble(x = seq(0, 15, 0.01), y = likelihoods)
p1 <- ggplot(df1, aes(x = x, y = y)) + geom_line() + xlab(expression(lambda)) +
  ylab(expression(paste("p(y | ", lambda, ")"))) +
  geom_vline(xintercept = mean(Y.Sim), lty = 2) + theme_bw() +
  ggtitle("Sampling model")
p1 <- p1 + annotate("text", x = mean(Y.Sim)+0.8, y=max(likelihoods), label = "MLE")
df2 <- tibble(x = seq(0, 15, 0.01), y1 = dgamma(x=seq(0, 15, 0.01),
  shape = a_0, rate = b_0),
  y2 = dgamma(x=seq(0, 15, 0.01), shape = a_n, rate = b_n)) %>%
  pivot_longer(!x, names_to = "type", values_to = "prob")
p2 <- ggplot(df2, aes(x = x, y = prob, col = type, group = type)) + geom_line() +
  theme(legend.title = element_blank()) +
  ggtitle("Prior and Posterior") + theme_bw() +
  scale_fill_discrete(labels = c(
    expression(paste("Prior for ", lambda)),
    expression(paste("Posterior for ", lambda))
  )) + xlab(expression(lambda)) + ylab(expression(paste("p(", lambda, " | a, b)")))
p <- grid.arrange(p1, p2, ncol = 1)
```

## Sampling model




## Prior and Posterior



p

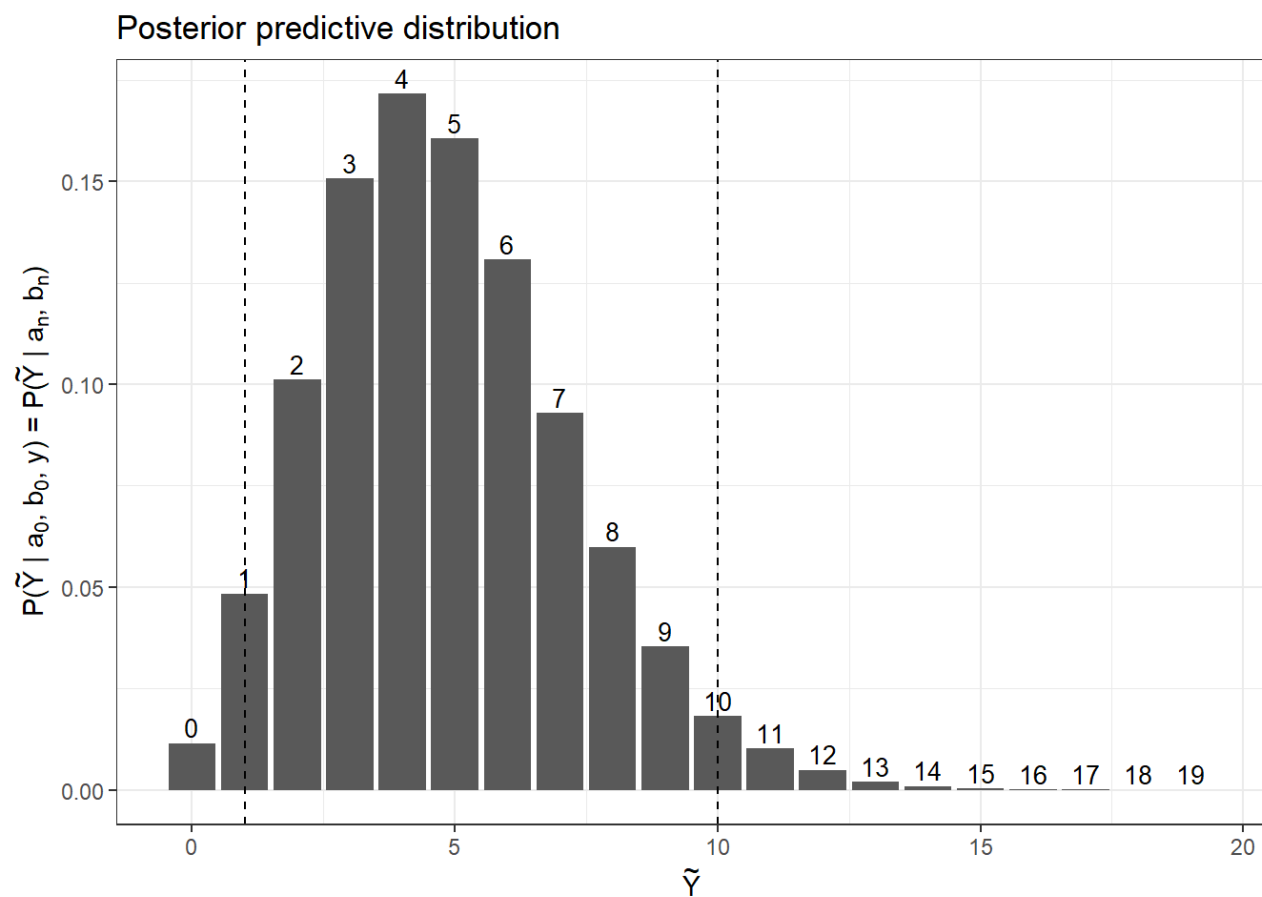
```
## TableGrob (2 x 1) "arrange": 2 grobs
##   z      cells   name      grob
## 1 1 (1-1,1-1) arrange gtable[layout]
## 2 2 (2-2,1-1) arrange gtable[layout]
```

```
# Illustration of the approximated joint posterior distribution of Lambda and the Poisson-distributed predictand
jointPostPlugInApprox = ks::kde(x = cbind(lambdaPostSim, Y.tildePostSim))
par(mfrow = c(1, 1))
plot(
  x=jointPostPlugInApprox, # For help plotting kde objects, see ?plot.kde
  cont = c(50, 95), # Percentiles to draw highest-probability regions
  xlim = c(2,8), ylim = c(0, 15),
  xlab = expression(lambda), ylab = expression(tilde(Y))
)
```



```
# Illustration of the approximated marginal posterior distribution of the Poisson-dist
ributed estimand
# 95% highest-probability predictive predictive interval
df <- tibble(x = as.integer(rownames(table(Y.tildePostSim))),
             y = table(Y.tildePostSim)/S)
ggplot(df, aes(x=x, y=y)) + geom_bar(stat="identity") +
  geom_text(aes(label=x), vjust=-0.3, size=3.5)+theme_minimal() +
  xlab(expression(tilde(Y))) + ylab(expression(paste("P(",tilde(Y)," | a"[0],", b"[0],
", y) = P(",tilde(Y)," | a"[n],", b"[n],")")))+
  geom_vline(xintercept = quantile(x = Y.tildePostSim, probs = c(0.025,.975)), lty = 2
) +
  ggtitle("Posterior predictive distribution")+ theme_bw()
```

```
## Don't know how to automatically pick scale for object of type table. Defaulting to
continuous.
```




## Comparing prior and posterior predictive distributions

```

# Illustration comparing the approximated marginal prior and posterior distributions of the Poisson-distributed predictand
check_occurences_1 <- function(x) {
  return(sum(Y.tildeSim %in% x))
}
check_occurences_2 <- function(x) {
  return(sum(Y.tildePostSim %in% x))
}
df <- tibble(x = 0:20, y1 = sapply(0:20, check_occurences_1)/S,
             y2 = sapply(0:20, check_occurences_2)/S) %>%
  pivot_longer(!x, names_to = "type", values_to = "probs")
ggplot(df) + geom_bar(stat='identity',
                     aes(x = x, y = probs, fill = factor(type)),
                     position=position_dodge(width = .7), width = .7)+
  theme_bw() + xlab(expression(tilde(Y))) +
  ylab(expression(paste("P(",tilde(Y)," | a, b)"))) +
  geom_vline(xintercept = quantile(x = Y.tildeSim, probs = c(0.025,0.975)),
             lty = 2, col = "black") +
  geom_vline(xintercept = quantile(x = Y.tildePostSim, probs = c(0.025,0.975)),
             lty = 2, col = "dark gray") +
  theme(legend.title = element_blank()) +
  scale_fill_discrete(labels = c("Prior predictive distribution",
                                "Posterior predictive distribution"))

```



# Simulating posterior (or prior) predictive distributions for whole samples

- Rather than drawing a single observation from the DGP model for each simulated value of  $\theta_m$ , we can draw whole samples if we wish:

$$\begin{aligned} \{\tilde{Y}_1^{(1)}, \dots, \tilde{Y}_n^{(1)}\} &\stackrel{iid}{\sim} \text{Dist}_m(\theta_m^{(1)}) \\ \{\tilde{Y}_1^{(2)}, \dots, \tilde{Y}_n^{(2)}\} &\stackrel{iid}{\sim} \text{Dist}_m(\theta_m^{(2)}) \\ &\vdots \\ \{\tilde{Y}_1^{(S)}, \dots, \tilde{Y}_n^{(S)}\} &\stackrel{iid}{\sim} \text{Dist}_m(\theta_m^{(S)}) \end{aligned}$$

- If (A) a particular summary statistic is calculated for each sample and (B) each sample is identical in size to the original sample, this simulation yields an approximation of a posterior predictive distribution of the chosen summary statistic, providing a Bayesian equivalent to a frequentist sampling distribution.

```
# Simulating a posterior predictive distribution of sample means for a Poisson model
Sample.tildePostSim = apply(
  X = as.matrix(lambdaPostSim, ncol=1), MARGIN = 1, FUN = rpois, n=5
)
# Plotting the predictive distribution and comparing to the observed sample mean
plot(
  x=table(colMeans(Sample.tildePostSim))/S,
  xlab = expression(tilde(bar(Y))),
  ylab = expression(paste("P(", tilde(bar(Y)), " | a"[n], ", b"[n], ")")),
  ylim = c(0, 0.1)
)
abline(v = quantile(colMeans(Sample.tildePostSim), probs = c(0.025, 0.975)), lty = 2)
abline(v = mean(Y.Sim), lty = 2, col = "blue", lwd = 2)
legend(
  "topright",
  legend = c(
    "Posterior predictive distribution of sample means",
    "95% posterior predictive interval",
    "Observed sample average"
  ),
  lwd = c(2, 1, 2), lty = c(1, 2, 2), col = c("black", "black", "blue")
)
```

