

# Lab 4

Alex Ziyu Jiang

(The code materials are kindly provided by Professor Carlos Cinelli.)

## Bayesian multiple linear regression: Saratoga Housing example

Today we look at an example of Bayesian multiple linear regression and show how to implement it using sampling softwares such as jags and rstan. Moreover, under a specific setting we show the relationship between Bayesian multiple linear regression and ridge regression, a frequentist regularization method commonly used in machine learning.

### model prerequisites

As usual, we load the package we need (remember to install them first using `install.packages()` if you haven't done so:

```
rm(list = ls())  
library(rjags)
```

```
## Loading required package: coda
```

```
## Linked to JAGS 4.3.0
```

```
## Loaded modules: basemod,bugs
```

```
library(rstan)
```

```
## Loading required package: StanHeaders
```

```
## Loading required package: ggplot2
```

```
## rstan (Version 2.21.3, GitRev: 2e1f913d3ca3)
```

```
## For execution on a local, multicore CPU with excess RAM we recommend calling
```

```
## options(mc.cores = parallel::detectCores()).
```

```
## To avoid recompilation of unchanged Stan programs, we recommend calling
```

```
## rstan_options(auto_write = TRUE)
```

```
##
```

```
## Attaching package: 'rstan'
```

```
## The following object is masked from 'package:coda':
```

```
##
```

```
##      traceplot
```

```
library(glmnet)
```

```
## Warning: package 'glmnet' was built under R version 4.1.2
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 4.1-4
```

```

library(rethinking)

## Loading required package: cmdstanr
## This is cmdstanr version 0.5.1.9000
## - CmdStanR documentation and vignettes: mc-stan.org/cmdstanr
## - CmdStan path: /Users/alexziyujiang/.cmdstan/cmdstan-2.29.2
## - CmdStan version: 2.29.2
## Loading required package: parallel
## rethinking (Version 2.21)
##
## Attaching package: 'rethinking'
## The following object is masked from 'package:rstan':
##
##      stan
## The following object is masked from 'package:stats':
##
##      rstudent
options(scipen = 99)

```

## data preprocessing

Then we load in the dataset and do some preprocessing to make the data suitable for our analysis. Just as a recap: for Bayesian regression models, standardizing (centering and rescaling by standard deviation) the predictors will lead to better MCMC sampling efficiency (Markov chains converge to equilibrium quicker) and easier prior choices, so here we standardize the columns in the data matrix.

```

# Load data -----
df <- read.csv("SaratogaHouses.csv")

# create matrix
x <- model.matrix(price ~ ., data = df)

# scale variables (except constant)
x.scale <- x
x.scale[,-1] <- apply(x.scale[,-1], 2, function(z) (z - mean(z))/sd(z))
y.scale <- df$price/sd(df$price)

# save sd's to rescale back
sdx <- apply(x[,-1], 2, sd)
sdy <- sd(df$price)

# function to rescale coefficients
rescale <- function(beta){
  beta <- beta*sdy/sdx
  names(beta) <- colnames(x)[-1]
  beta
}

```

## model form

Recall that here we are building a multivariate linear regression model. If the model matrix is  $\mathbf{X}$ , the vector of regression coefficients are  $\beta$ , the model has the following form:

$$\begin{aligned}y_i &\sim \text{Normal}(\mu_i, \sigma^2), i = 1, \dots, n \\ \mu &= \mathbf{X}\beta \\ \beta_0 &\sim \text{Normal}(\mu_0, s_0^2) \\ \beta_j &\sim \text{Normal}(\mu_\beta, s_\beta^2), j = 1, \dots, p \\ \sigma &\sim \text{Exponential}(\lambda)\end{aligned}$$

Here  $\mu_\beta$  and  $s_\beta^2$  represents the mean and variance for the regression coefficients  $\beta_j$ , and  $\lambda$  is the rate coefficient of the exponential prior for  $\sigma$ . Note that  $\mu_\beta$ ,  $s_\beta^2$  and  $\lambda$  are ‘parameters’ of the prior distributions, so instead of placing prior on them we feed them actual values – these parameters are called **hyperparameters** in Bayesian statistics. We choose  $\lambda = 1$ , a relatively ‘flat’ prior for the standard deviation. For the regression coefficients, we choose a bunch of normal priors with mean zero, because we don’t really have prior knowledge about how each effect will look like before fitting the data. For the intercept model, we place a flat normal prior with standard deviation  $s_0 = 1,000$ . For the other variables, we place a ‘tighter’ regularization prior on all of these variables with standard deviation  $s_\beta = 0.02$ .

## model implementation

### JAGS

Finally, we compile the JAGS and STAN code for model implementation (note the notational difference regarding the normal variance/precision):

```
# JAGS -----

# generic model code
linear_model_code <- "
  data{
    D <- dim(x)
    n <- D[1]
    p <- D[2]
  }
  model{
    for(i in 1:n){
      # likelihood
      y[i] ~ dnorm(mu[i], tau)
      # # posterior predictive
      # ynew[i] ~ dnorm(mu[i], tau)
    }
    # conditional mean using matrix algebra
    mu <- x %*% beta
    for(j in 1:p){
      beta[j] ~ dnorm(mb[j], pow(sb[j], -2))
    }
    sigma ~ dexp(lambda)
    tau <- pow(sigma, -2)
  }
"

# flat prior for constant
```

```

# tight regularizing priors for all other parameters
model <- jags.model(file = textConnection(linear_model_code),
                    data = list(x = x.scale,
                                y = y.scale,
                                mb = rep(0, ncol(x)),
                                sb = c(1000, rep(.02, ncol(x)-1)),
                                lambda = 1))

```

```

## Compiling data graph
##   Resolving undeclared variables
##   Allocating nodes
##   Initializing
##   Reading data back into data table
## Compiling model graph
##   Resolving undeclared variables
##   Allocating nodes
## Graph information:
##   Observed stochastic nodes: 1728
##   Unobserved stochastic nodes: 20
##   Total graph size: 36359
##
## Initializing model

```

```

nsim <- 5e3
# burn in
update(model, n.iter = nsim)

# samples
samps <- coda.samples(model = model, n.iter = nsim,
                      variable.names = c("beta", "sigma"))

# check trace plots
# plot(samps)

# transform back to original scale
samps.df <- as.data.frame(samps[[1]])
post.means <- apply(samps.df, 2, mean)[-c(1,20)]
post.means <- rescale(post.means)
post.means

```

##	lotSize	age	landValue
##	7432.0780941	-113.1577644	0.6412473
##	livingArea	pctCollege	bedrooms
##	35.5364074	144.2256107	3464.7609200
##	fireplaces	bathrooms	rooms
##	8884.6695971	22264.5847407	4423.9892157
##	heatinghot air	heatinghot water/steam	fuelgas
##	6488.3455481	-1972.1805342	5584.9854936
##	fueloil	sewerpublic/commercial	sewerseptic
##	-100.4664115	184.6637635	-436.5662278
##	waterfrontYes	newConstructionYes	centralAirYes
##	94370.7441953	-10778.6115139	12958.2920648

We run the model and generate 5,000 posterior samples for  $\beta$  and  $\sigma$ . We can use the posterior samples for  $\beta$  to calculate its posterior mean. Finally we transform it back to the original scale for clearer interpretation (in the sense that ‘the rate of change’ is associated with unit change in the actual variables).

## STAN

Similarly, we could do the stan version:

```
## Running MCMC with 1 chain...
##
## Chain 1 Iteration:    1 / 5000 [  0%] (Warmup)
## Chain 1 Iteration:   100 / 5000 [  2%] (Warmup)

## Chain 1 Informational Message: The current Metropolis proposal is about to be rejected because of the
## Chain 1 Exception: normal_lpdf: Scale parameter is 0, but must be positive! (in '/var/folders/42/k2f
## Chain 1 If this warning occurs sporadically, such as for highly constrained variable types like covar
## Chain 1 but if this warning occurs often then your model may be either severely ill-conditioned or m
## Chain 1

## Chain 1 Iteration:   200 / 5000 [  4%] (Warmup)
## Chain 1 Iteration:   300 / 5000 [  6%] (Warmup)
## Chain 1 Iteration:   400 / 5000 [  8%] (Warmup)
## Chain 1 Iteration:   500 / 5000 [ 10%] (Warmup)
## Chain 1 Iteration:   600 / 5000 [ 12%] (Warmup)
## Chain 1 Iteration:   700 / 5000 [ 14%] (Warmup)
## Chain 1 Iteration:   800 / 5000 [ 16%] (Warmup)
## Chain 1 Iteration:   900 / 5000 [ 18%] (Warmup)
## Chain 1 Iteration:  1000 / 5000 [ 20%] (Warmup)
## Chain 1 Iteration:  1100 / 5000 [ 22%] (Warmup)
## Chain 1 Iteration:  1200 / 5000 [ 24%] (Warmup)
## Chain 1 Iteration:  1300 / 5000 [ 26%] (Warmup)
## Chain 1 Iteration:  1400 / 5000 [ 28%] (Warmup)
## Chain 1 Iteration:  1500 / 5000 [ 30%] (Warmup)
## Chain 1 Iteration:  1600 / 5000 [ 32%] (Warmup)
## Chain 1 Iteration:  1700 / 5000 [ 34%] (Warmup)
## Chain 1 Iteration:  1800 / 5000 [ 36%] (Warmup)
## Chain 1 Iteration:  1900 / 5000 [ 38%] (Warmup)
## Chain 1 Iteration:  2000 / 5000 [ 40%] (Warmup)
## Chain 1 Iteration:  2100 / 5000 [ 42%] (Warmup)
## Chain 1 Iteration:  2200 / 5000 [ 44%] (Warmup)
## Chain 1 Iteration:  2300 / 5000 [ 46%] (Warmup)
## Chain 1 Iteration:  2400 / 5000 [ 48%] (Warmup)
## Chain 1 Iteration:  2500 / 5000 [ 50%] (Warmup)
## Chain 1 Iteration:  2501 / 5000 [ 50%] (Sampling)
## Chain 1 Iteration:  2600 / 5000 [ 52%] (Sampling)
## Chain 1 Iteration:  2700 / 5000 [ 54%] (Sampling)
## Chain 1 Iteration:  2800 / 5000 [ 56%] (Sampling)
## Chain 1 Iteration:  2900 / 5000 [ 58%] (Sampling)
## Chain 1 Iteration:  3000 / 5000 [ 60%] (Sampling)
## Chain 1 Iteration:  3100 / 5000 [ 62%] (Sampling)
## Chain 1 Iteration:  3200 / 5000 [ 64%] (Sampling)
## Chain 1 Iteration:  3300 / 5000 [ 66%] (Sampling)
## Chain 1 Iteration:  3400 / 5000 [ 68%] (Sampling)
## Chain 1 Iteration:  3500 / 5000 [ 70%] (Sampling)
## Chain 1 Iteration:  3600 / 5000 [ 72%] (Sampling)
## Chain 1 Iteration:  3700 / 5000 [ 74%] (Sampling)
## Chain 1 Iteration:  3800 / 5000 [ 76%] (Sampling)
## Chain 1 Iteration:  3900 / 5000 [ 78%] (Sampling)
```

```
## Chain 1 Iteration: 4000 / 5000 [ 80%] (Sampling)
## Chain 1 Iteration: 4100 / 5000 [ 82%] (Sampling)
## Chain 1 Iteration: 4200 / 5000 [ 84%] (Sampling)
## Chain 1 Iteration: 4300 / 5000 [ 86%] (Sampling)
## Chain 1 Iteration: 4400 / 5000 [ 88%] (Sampling)
## Chain 1 Iteration: 4500 / 5000 [ 90%] (Sampling)
## Chain 1 Iteration: 4600 / 5000 [ 92%] (Sampling)
## Chain 1 Iteration: 4700 / 5000 [ 94%] (Sampling)
## Chain 1 Iteration: 4800 / 5000 [ 96%] (Sampling)
## Chain 1 Iteration: 4900 / 5000 [ 98%] (Sampling)
## Chain 1 Iteration: 5000 / 5000 [100%] (Sampling)
## Chain 1 finished in 3.8 seconds.
```

```
stan.means <- apply(as.data.frame(m.stan), 2, mean)[-c(1,20,21)]
stan.means <- rescale(stan.means)
stan.means
```

```
##          lotSize          age          landValue
##      7382.6122462      -111.6198372      0.6423868
##      livingArea      pctCollege      bedrooms
##      35.5053727      143.8605749      3559.4846901
##      fireplaces      bathrooms      rooms
##      8906.9120395      22184.2658244      4409.4458256
##      heatinghot air heatinghot water/steam      fuelgas
##      6406.9875993      -2049.3089260      5656.3703135
##      fueloil sewerpublic/commercial      sewerseptic
##      -58.0251741      67.9907910      -487.1475680
##      waterfrontYes      newConstructionYes      centralAirYes
##      94308.0033961      -10502.0987866      12961.8292228
```

## QUAP() from the textbook

Finally, the textbook has a fancy function called `quap()`, to approximate the posterior distributions under regression settings. We repeat the similar analysis:

```
# Quadratic Approximation -----
# using quadratic approximation (your book)
model.quap <- quap(flist = alist(
  y ~ dnorm(mu, sigma),
  mu <- alpha + x %*% beta,
  alpha ~ dnorm(0, 1000),
  beta ~ dnorm(0, 0.02),
  sigma ~ dexp(1)),
  data = list(x = x.scale[, -1],
              y = y.scale),
  start = list(beta = rep(0, ncol(x) - 1)))
# transform back
quap.coef <- coef(model.quap)[-c(19, 20)]
quap.coef <- rescale(quap.coef)
quap.coef
```

```
##          lotSize          age          landValue
##      7402.5861991      -112.6974093      0.6430178
##      livingArea      pctCollege      bedrooms
##      35.6413884      144.3355290      3462.2684098
##      fireplaces      bathrooms      rooms
```

##	8889.9119816	22212.2662579	4422.3961462
##	heatinghot air	heatinghot water/steam	fuelgas
##	6454.1508202	-2020.5862189	5663.4477399
##	fueloil	sewerpublic/commercial	sewerseptic
##	-72.3232879	108.4333332	-484.0987777
##	waterfrontYes	newConstructionYes	centralAirYes
##	94543.4351029	-10759.8934385	12941.0703763

## Ridge regression and Bayesian linear regression

Let's think about frequentist linear regression model for a moment. The ordinary least squares estimate of a linear model can be reframed as an optimization problem:

$$\hat{\beta} = \underset{\beta}{\operatorname{argmin}} (y - X\beta)^T (y - X\beta).$$

As a remedy to model overfitting problems, ridge regression is a commonly used regularization method that tends to reduce the magnitude of each predictor variable in the model. To do ridge regression, we simply add an extra penalty term  $\lambda \|\beta\|_2^2$  that penalizes the Euclidean norm of the vector of coefficient. Here  $\lambda$  is a hyperparameter (a different 'hyperparameter' than the one in Bayesian statistics as there is no prior here) that controls how much you want to penalize the vector of coefficients.

$$\hat{\beta} = \underset{\beta}{\operatorname{argmin}} (y - X\beta)^T (y - X\beta) + \lambda \|\beta\|_2^2,$$

There is a Bayesian interpretation to the ridge regression framework: for a Bayesian linear regression model with fixed residual variance  $\sigma^2$  and independent Gaussian prior  $\mathcal{N}(0, \tau^2 \mathbf{I}_q)$  on the regression coefficients  $\beta$ , the posterior mode for  $\beta$ ,  $p(\beta \mid \mathbf{X}, \sigma^2, \tau^2)$  corresponds to the ridge regression estimate with  $\lambda = \frac{\sigma^2}{\tau^2}$ . We won't get into the reasoning behind this, but you can refer to here for more detail: <https://statisticaloddsandends.wordpress.com/2018/12/29/bayesian-interpretation-of-ridge-regression/>.

We will first fit the ridge regression estimates using functions in package `glmnet`. A couple of things to notice:

- 'alpha = 0' means we are doing the ridge regression (unrelated to today's material, but if you set alpha to be 1, we get lasso instead).
- Since the model we are considering is a little different from the setting above, we will not be getting exactly the same estimates (also, we are actually using posterior mean instead of posterior mode), but they should be similar.
- We will also be computing estimates without penalizing to show the difference between these estimates.

```
# OLS and Ridge -----

# fit Ridge for comparison
gl.out <- glmnet(x = x.scale[, -1], y = y.scale, standardize = F, intercept = T, alpha = 0, lambda = 0.4)
gl.coef <- coef(gl.out)[-1]

# transform back to original scale
gl.coef <- rescale(gl.coef)
gl.coef

##          lotSize          age          landValue
## 7599.0529985    -114.5099914      0.6813298
##    livingArea    pctCollege      bedrooms
## 38.0590035     121.5486238     2686.8866930
##    fireplaces      bathrooms      rooms
```

```
##           8408.1669380           22905.8049761           4431.8754235
## heatinghot air heatinghot water/steam           fuelgas
##           6430.8509086           -2289.0060878           5520.2448287
##           fueloil sewerpublic/commercial           sewerseptic
##           71.6557607           -9.8817011           -360.9500877
## waterfrontYes newConstructionYes           centralAirYes
##           99426.1087329           -13930.7062043           12915.6685082
```

```
# fit lm for comparison
lm.coef <- coef(lm(price ~ ., data = df))[-1]
```

```
# compare estimates
round(cbind(`lm (not regularized)` = lm.coef,
            jags = post.means,
            stan = stan.means,
            quap = quap.coef,
            glmnet = gl.coef),3)
```

	lm (not regularized)	jags	stan	quap
## lotSize	7599.449	7432.078	7382.612	7402.586
## age	-130.446	-113.158	-111.620	-112.697
## landValue	0.922	0.641	0.642	0.643
## livingArea	69.960	35.536	35.505	35.641
## pctCollege	-110.159	144.226	143.861	144.336
## bedrooms	-7835.192	3464.761	3559.485	3462.268
## fireplaces	1036.613	8884.670	8906.912	8889.912
## bathrooms	23112.452	22264.585	22184.266	22212.266
## rooms	3019.761	4423.989	4409.446	4422.396
## heatinghot air	82.452	6488.346	6406.988	6454.151
## heatinghot water/steam	-10372.246	-1972.181	-2049.309	-2020.586
## fuelgas	10931.274	5584.985	5656.370	5663.448
## fueloil	6550.471	-100.466	-58.025	-72.323
## sewerpublic/commercial	3321.168	184.664	67.991	108.433
## sewerseptic	4845.107	-436.566	-487.148	-484.099
## waterfrontYes	120193.978	94370.744	94308.003	94543.435
## newConstructionYes	-45443.421	-10778.612	-10502.099	-10759.893
## centralAirYes	9953.091	12958.292	12961.829	12941.070
## glmnet				
## lotSize	7599.053			
## age	-114.510			
## landValue	0.681			
## livingArea	38.059			
## pctCollege	121.549			
## bedrooms	2686.887			
## fireplaces	8408.167			
## bathrooms	22905.805			
## rooms	4431.875			
## heatinghot air	6430.851			
## heatinghot water/steam	-2289.006			
## fuelgas	5520.245			
## fueloil	71.656			
## sewerpublic/commercial	-9.882			
## sewerseptic	-360.950			
## waterfrontYes	99426.109			
## newConstructionYes	-13930.706			



## centralAirYes

12915.669

## Conclusion

- Bayesian linear regression can be viewed as a regularization method
- Some concluding remarks on covariate standardizing: centering and rescaling (1) helps sampling and (2) helps choosing priors:
  - In ridge regression we standardize the covariates and give them the same ‘penalty term’, for the Bayesian equivalent, instead of doing the penalty term we place a tight prior with large precision around zero for all of the variables
  - For the ‘common penalty’ (in terms of the tight prior) to make sense, we need to do the standardization